NORTHWESTERN UNIVERSITY | SCHOOL OF PROFESSIONAL STUDIES

PREDICT 422: Practical Machine Learning **Charity Project: Part 3**

Problem Description

A charitable organization wishes to develop a machine learning model to improve the cost-effectiveness of their direct marketing campaigns to previous donors. According to their recent mailing records, the typical overall response rate is 5.1%. Out of those who respond (donate) to the mailing, the average donation is $15.62. Each mailing costs $0.99 to produce and send. The mailing includes a gift of personalized address labels and an assortment of cards and envelopes. It is not cost-effective to mail everyone because the expected profit from each mailing is $15.62 x 0.051 – $0.99 = - $0.19.

We will address this problem over a series of assignments. The overall goal for this problem is to maximize the net profit of the next direct marketing campaign. Our approach will be two-fold:

1. We would like to build a **regression** model to predict expected gift amounts from donors.
2. We would like to develop a **classification** model that can effectively capture likely donors.

The overall problem will be broken down into four separate assignments.

1. Exploratory Data Analysis
2. The Regression Problem
3. **The Classification Problem**
4. The Mailing List Problem

Charity Problem — Part 2

Data Files

- `dataDict.txt`
- `trainSample.csv`

Sample Code

- `SampleCodePart3.R`

Exercises

1. Import Data

    a. Read the data into R from the CSV file `trainSample.csv` the same as you did in Part 1 Exercise 1.
    b. For this assignment (Part 3, The Classification Problem) you will use the full dataset (no subsetting). Name the dataset `classData` for the purposes of this assignment.

2. Data Preparation

    a. There are two options here. Either repeat the data processing steps you used in Part 2 to the Part 3 dataset `classData` OR make modifications or additions to the Part 2 data processing for the purposes of Part 3.
    b. Briefly describe the data preparation steps that you take. In particular, make note of any differences (if any) between the data processing in Part 2 and Part 3 of the project.

3. Dataset Partitioning

   For this assignment, you will employ a hold-out test dataset for model validation and selection.

   a. Hold-Out Test Set
      The first step you should take is to sample 25% of the observations in the dataset to form a hold-out test set. This data will be referred to as the **Classification Test Set** (or simply the Test Set for the remainder of this document). Report the number of observations and the distribution of response values in the Test Set. The data in the Test Set should not be used until Exercise 5 of this assignment.
   b. Training Set
      The remaining 75% of the observations will be referred to as the **Classification Training Set** (or simply the Training Set for the remainder of this document). Report the number of observations and the distribution of response values in the Training Set.

4. Model Fitting

   Use R to develop various models for the response variable DONR. The variables ID and DAMT are not to be used as predictors. Fit at least one model from each of the following four categories. Each model should be fit to the Training Set data only.

   a. Simple logistic regression (ISLR Chapter 4) [Recall that simple logistic regression is logistic regression with a *single predictor variable*.]
   b. Multiple logistic regression or Linear Discriminant Analysis (ISLR Chapter 4)
   c. Tree-based models (ISLR Chapter 8)
   d. Another model of your choice, which may include a second model from one of the three prior categories

   For each model, report the form of the model you are fitting (e.g. the formula used to specify the model). Explain the reasoning for why you are fitting a model of that form (e.g. for simple logistic regression, explain how you selected which predictor to use). Explain any hyper-parameter tuning that you do (e.g. tuning the threshold value for logistic regression). Report summary and diagnostic information as appropriate for each model (in particular, confusion matrices and TP and FP rates).

5. Model Validation

   Use R to perform model validation on the models you fit in Exercise 4. The model validation process is outlined below.

   a. Build a table (in your document) that has one row for each model you fit in Exercise 4. The table should have seven columns (at minimum): Model Name, Training Set Accuracy, Training Set TP Rate, Training Set FP Rate, Test Set Accuracy, Test Set TP Rate, and Test Set FP Rate. You can include additional columns if you would like.
   b. For the Training Set MSE, predict DONR for all of the individuals in the Training Set, and calculate the MSE from the Training Set predictions. Note that it is expected that this MSE value will *underestimate* the test error. The Training Set MSE is included in the Model Validation table for comparison purposes only. If there is a dramatic difference between the Training Set MSE and the Test Set MSE, then that is an indication that the model has overfit the training data.
   c. For the Test Set MSE, predict DONR for all of the individuals in the Test Set, and calculate the MSE from the Test Set predictions.. Note that you do not retrain or refit the model to the Test Set data, nor do you re-tune the hyper-parameters.
   d. In addition, present the Test Set lift chart for each model that you build.

6. Model Selection

   Use the table you generated in Exercise 5 to select the best model to carry forward to Part 4 of the Charity Project.

   a. Comment on the predictive accuracy you get from your models.
   b. Explain which of your models you select as being the best performing model and why. Note that model selection should be based on the Test Set metric values. If two models have similar Test Set metric values, then the model with fewer predictors should be selected.

Submissions

Submit the following files in Canvas:

1. PDF or Word document that details your findings from the exercises. Include figures and tables as applicable. Clearly indicate the exercise number in your document.
2. Your R code (if more than one .r or .R file, zip them into a single file for upload).